

```
import pandas as pd
import pylab
import numpy as np
import sklearn
from sklearn import linear_model
import sklearn.preprocessing as preprocessing
import sklearn.metrics as metrics
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
from statsmodels.stats import proportion
```

▼ Load the Dataset

Data set can be found from this link: <https://archive.ics.uci.edu/ml/datasets/student+performance>.

```
!wget https://archive.ics.uci.edu/ml/machine-learning-databases/00320/student.zip
```

```
!rm ./*.csv
!rm ./*.R
!rm student.txt
!unzip student.zip
```

```
features = ["school", "sex", "age", "address", "famsize", "Pstatus",
            "Medu", "Fedu", "Mjob", "Fjob", "reason", "guardian",
            "traveltime", "studytime", "failures", "schoolsup", "famsup", "paid",
            "activities", "nursery", "higher", "internet", "romantic", "famrel",
            "freetime", "goout", "Dalc", "Walc", "health", "absences", "G1", "G2", "G3"]
```

```
df = pd.read_csv('student-por.csv', sep=';', skiprows=1, engine='python', names=features, na_v
```

```
df
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	reason
0	GP	F	18	U	GT3	A	4	4	at_home	teacher	co

1	GP	F	17	U	GT3	T	1	1	at_home	other	co
2	GP	F	15	U	LE3	T	1	1	at_home	other	c
3	GP	F	15	U	GT3	T	4	2	health	services	h
4	GP	F	16	U	GT3	T	3	3	other	other	h
...
644	MS	F	19	R	GT3	T	2	3	services	other	co
645	MS	F	18	U	LE3	T	3	1	teacher	services	co

```
print('Total number of participants')
print(len(df))
print( )
```

```
female = df[df['sex'] == "F"]
print('Number of female participants:')
print(len(female))
print( )
```

```
male = df[df['sex'] == "M"]
print('Number of male participants:')
print(len(male))
print( )
```

```
print('Percentage of female participants:')
print(len(female)/len(df))
print( )
```

```
print('Percentage of female participants:')
print(len(male)/len(df))
print( )
```

Total number of participants
649

Number of female participants:
383

Number of male participants:
266

Percentage of female participants:
0.5901386748844376

Percentage of female participants:

0.4098613251155624

```
print('Total number of participants')
```

```
print(len(df))
print( )

urban = df[df['address'] == "U"]
print('Number of urban participants:')
print(len(urban))
print( )

rural = df[df['address'] == "R"]
print('Number of rural participants:')
print(len(rural))
print( )

print('Percentage of urban participants:')
print(len(urban)/len(df))
print( )

print('Percentage of rural participants:')
print(len(rural)/len(df))
print( )

urban_succeed = urban[urban['G3'] > 9]
urban_succeed

urban = df[df['address'] == "U"]
print('Number of urban students succeeding:')
print(len(urban_succeed))
print( )

print('Number of urban students failed:')
print(len(urban) - len(urban_succeed))
print( )

print('Percent of urban students succeeding:')
print(len(urban_succeed)/len(urban))
print( )

rural_succeed = rural[rural['G3'] > 9]
rural_succeed

rural = df[df['address'] == "R"]
print('Number of rural students succeeding:')
print(len(rural_succeed))
print( )

print('Number of rural students failed:')
print(len(rural) - len(rural_succeed))
print( )

print('Percent of rural students succeeding:')
print(len(rural_succeed)/len(rural))
print( )
```

```
print( )
```

```
Total number of participants  
649
```

```
Number of urban participants:  
452
```

```
Number of rural participants:  
197
```

```
Percentage of urban participants:  
0.6964560862865947
```

```
Percentage of rural participants:  
0.3035439137134052
```

```
Number of urban students succeeding:  
396
```

```
Number of urban students failed:  
56
```

```
Percent of urban students succeeding:  
0.8761061946902655
```

```
Number of rural students succeeding:  
153
```

```
Number of rural students failed:  
44
```

```
Percent of rural students succeeding:  
0.7766497461928934
```

```
female_succeed = female[female['G3'] > 9]  
female_succeed
```

```
female = df[df['sex'] == "F"]  
print('Number of female students succeeding:')  
print(len(female_succeed))  
print( )
```

```
print('Number of female students failed:')  
print(len(female) - len(female_succeed))  
print( )
```

```
print('Percent of female students succeeding:')
```

```
print(len(female_succeed)/len(female))  
print( )
```

```
Number of female students succeeding:
```

```
Number of female students succeeding:
333
```

```
Number of female students failed:
50
```

```
Percent of female students succeeding:
0.8694516971279374
```

```
male_succeed = male[male['G3'] > 9]
male_succeed
```

```
male = df[df['sex'] == "M"]
print('Number of male students succeeding:')
print(len(male_succeed))
print( )
```

```
print('Number of male students failed:')
print(len(male) - len(male_succeed))
print( )
```

```
print('Percent of male students succeeding:')
print(len(male_succeed)/len(male))
print( )
```

```
Number of male students succeeding:
216
```

```
Number of male students failed:
50
```

```
Percent of male students succeeding:
0.8120300751879699
```

▼ Cleaning the Data

```
df.head()
labels = df['G3']
labels = labels.replace(1, 0).replace(2, 0).replace(3, 0).replace(4, 0)
labels = labels.replace(5, 0).replace(6, 0).replace(7, 0).replace(8, 0)
labels = labels.replace(9, 0).replace(10, 1)
labels = labels.replace(11, 1).replace(12, 1).replace(13, 1).replace(14, 1)
labels = labels.replace(15, 1).replace(16, 1).replace(17, 1).replace(18, 1)
labels = labels.replace(19, 1).replace(20, 1)

print(labels)
```

```

0      1
1      1
2      1
3      1
4      1
..
644    1
645    1
646    0
647    1
648    1
Name: G3, Length: 649, dtype: int64

```

```
data = df.drop(columns=["G3"])
```

```

#0 for Gabriel Pereira, 1 for Mousinho da Silveira
school = {'GP': 0, 'MS': 1}
data.school = [school[item] for item in data.school]

```

```

#0 for female, 1 for male
sex = {'F': 0, 'M': 1}
data.sex = [sex[item] for item in data.sex]

```

```

#0 for urban, 1 for rural
address = {'U': 0, 'R': 1}
data.address = [address[item] for item in data.address]

```

```

#0 for less than or equal to 3, 1 for greater than three
famsize = {'LE3': 0, 'GT3': 1}
data.famsize = [famsize[item] for item in data.famsize]

```

```

#0 for living together, 1 for living apart
Pstatus = {'T': 0, 'A': 1}
data.Pstatus = [Pstatus[item] for item in data.Pstatus]

```

```

#0 for teacher, 1 for health care, 2 for services, 3 for at home, 4 for other
Mjob = {'teacher': 0, 'health': 1, 'services': 2, 'at_home': 3, 'other': 4}
data.Mjob = [Mjob[item] for item in data.Mjob]
data.Fjob = [Mjob[item] for item in data.Fjob]

```

```

#0 for close to home, 1 for school reputation, 2 for course preference, 3 for other
reason = {'home': 0, 'reputation': 1, 'course': 2, 'other': 3}
data.reason = [reason[item] for item in data.reason]

```

```

#0 for mother, 1 for father, 2 for other
guardian = {'mother': 0, 'father': 1, 'other': 2}
data.guardian = [guardian[item] for item in data.guardian]

```

```

#0 for yes, 1 for no
yesOrNo = {'yes': 0, 'no': 1}
data.schoolsup = [yesOrNo[item] for item in data.schoolsup]
data.famsup = [yesOrNo[item] for item in data.famsup]
data.reason = [yesOrNo[item] for item in data.reason]

```

```

data.paid = [yesOrNo[item] for item in data.paid]
data.activities = [yesOrNo[item] for item in data.activities]
data.nursery = [yesOrNo[item] for item in data.nursery]
data.higher = [yesOrNo[item] for item in data.higher]
data.internet = [yesOrNo[item] for item in data.internet]
data.romantic = [yesOrNo[item] for item in data.romantic]

train_proportion = 0.7
train_x, test_x, train_y, test_y = train_test_split(data, labels,
                                                    train_size=train_proportion,
                                                    test_size=1-train_proportion,
                                                    random_state=419)

```

data

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	reason	
0	0	0	18	0	1	1	4	4	3	0	2	
1	0	0	17	0	1	0	1	1	3	4	2	
2	0	0	15	0	0	0	1	1	3	4	3	
3	0	0	15	0	1	0	4	2	1	2	0	
4	0	0	16	0	1	0	3	3	4	4	0	
...	
644	1	0	19	1	1	0	2	3	2	4	2	
645	1	0	18	0	0	0	3	1	0	2	2	
646	1	0	18	0	1	0	1	1	4	4	2	
647	1	1	17	0	0	0	3	1	2	2	2	
648	1	1	18	1	0	0	3	2	2	4	2	

649 rows × 32 columns

▼ Data Exploration

Let's look at the data distribution between female and male students to better inform our analysis.

```

all_female = data[(data['sex'] == 0)]

all_male = data[(data['sex'] == 1)]
print(len(all_female) / len(data))
print(len(all_male) / len(data))

```

```
0.5901386748844376
0.4098613251155624
```

So, there's about a 60/40 split with more female students than male students.

Let's look at the distribution between urban and rural students.

```
all_urban = data[(data['address'] == 0)]
all_rural = data[(data['address'] == 1)]
print(len(all_urban) / len(data))
print(len(all_rural) / len(data))
```

```
0.6964560862865947
0.3035439137134052
```

We can see the majority of our original data favors urban students over rural students with about a 70/30 split.

Let's look at our train and test data sets as well

```
test_female = test_x[(test_x['sex'] == 0)]
test_male = test_x[(test_x['sex'] == 1)]
test_urban = test_x[(test_x['address'] == 0)]
test_rural = test_x[(test_x['address'] == 1)]

print("Test data")
print("-----")
print("Female students:", len(test_female) / len(test_x))
print("Male students: ", len(test_male) / len(test_x))
print()
print("Urban students:", len(test_urban) / len(test_x))
print("Rural students:", len(test_rural) / len(test_x))
```

```
Test data
-----
Female students: 0.5435897435897435
Male students:   0.4564102564102564

Urban students: 0.6461538461538462
Rural students: 0.35384615384615387
```

```
train_female = train_x[(train_x['sex'] == 0)]
train_male = train_x[(train_x['sex'] == 1)]
train_urban = train_x[(train_x['address'] == 0)]
train_rural = train_x[(train_x['address'] == 1)]

print("Train data")
print("-----")
```



```
print( ----- )
print("Female students:", len(train_female) / len(train_x))
print("Male students:  ", len(train_male) / len(train_x))
print()
print("Urban students:", len(train_urban) / len(train_x))
print("Rural students:", len(train_rural) / len(train_x))
```

Train data

Female students: 0.6101321585903083

Male students: 0.3898678414096916

Urban students: 0.7180616740088106

Rural students: 0.28193832599118945

▼ Independence

```
from sklearn import linear_model
```

```
num_train = len(train_x)
```

```
cls = linear_model.LogisticRegression(max_iter=100000)
```

```
cls.fit(train_x, train_y)
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100000,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                    warm_start=False)
```

```
predictions = cls.predict(test_x)
```

```
overall_scores = cls.score(test_x, test_y)
```

```
overall_scores
```

```
0.9230769230769231
```

```
from sklearn.metrics import classification_report
```

```
print(classification_report(predictions, test_y))
```

	precision	recall	f1-score	support
0	0.68	0.76	0.72	25
1	0.96	0.95	0.96	170
accuracy			0.92	195
macro avg	0.82	0.85	0.84	195
weighted avg	0.93	0.92	0.92	195

▼ Demographic Parity

```
scores = cls.predict_proba(test_x)[: , 1]
d = {'target' : test_y.values,
     'score' : scores,
     'prediction' : predictions,
     'address' : test_x['address'],
     'gender' : test_x['sex']}
```

```
marginals = pd.DataFrame(data=d, columns=['target', 'score', 'prediction', 'address', 'gender'])
marginals.head()
```



	target	score	prediction	address	gender
46	1	0.999177	1	0	0
197	1	0.999706	1	0	0
208	1	0.999889	1	0	1
326	1	0.999714	1	0	1
359	1	0.999999	1	1	0

▼ For Female vs. Male

```
postive_class = marginals[(marginals['prediction'] == 1) ]

positive_female = postive_class[(postive_class['gender'] == 0)]
positive_male = postive_class[(postive_class['gender'] == 1)]

print("# of positively classified males to females ->", len(positive_male)/len(positive_female))

female = marginals[marginals['gender'] == 0]
male = marginals[marginals['gender'] == 1]

print("% of positively classified males to females ->", (len(positive_male)/len(male))/(len(positive_female)/len(female)))

# of positively classified males to females -> 0.8085106382978723
% of positively classified males to females -> 0.9629452546019602
```

▼ For Urban vs. Rural

```
positive_urban = postive_class[(postive_class['address'] == 0)]
```

```

positive_rural = postive_class[(postive_class['address'] == 1)]

print("# of positively classified rural to urban students ->", len(positive_rural)/len(positi

urban = marginals[marginals['address'] == 0]
rural = marginals[marginals['address'] == 1]

print("% of positively classified rural to urban students ->", (len(positive_rural)/len(rural

# of positively classified rural to urban students -> 0.5315315315315315
% of positively classified rural to urban students -> 0.9706227967097533

```

We care more about the percentage comparison because

▼ Seperation

```

fpr_all, tpr_all, _ = metrics.roc_curve(marginals['target'], marginals['score'])

plt.figure(1, figsize = (6,5))
plt.title('ROC curve')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')

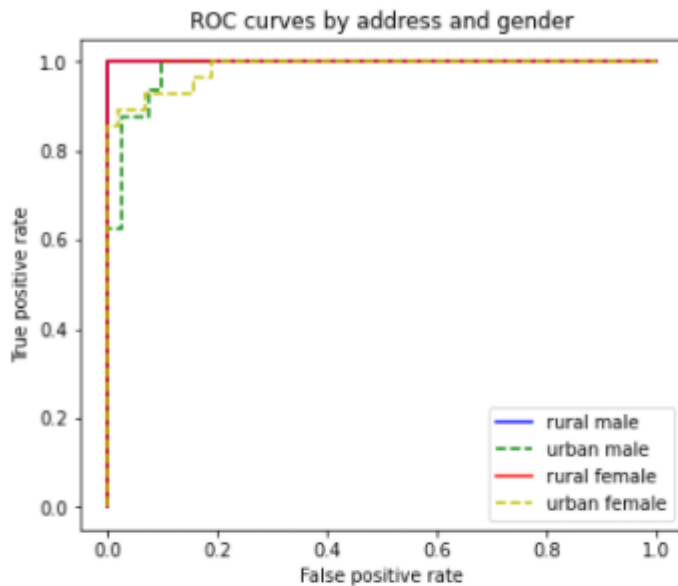
plt.plot(fpr_all, tpr_all, color='b')
plt.legend()
plt.show()

```

No handles with labels found to put in legend.


ROC curve

Please note that if a line does not show up on the ROC curve, it is most likely because none of the data points have a positive criteria or it overlaps with another ROC curve. This usually occurs with rural male because the red line takes precedence. Example Below:



```
plt.figure(1, figsize=(6, 5))
plt.title('ROC curves by address and gender')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')
plt.plot(fpr_ruralmale, tpr_ruralmale, label='rural male', line
```

[<matplotlib.lines.Line2D at 0x7fa8f77c8b10>]



```
male = marginals[marginals['gender'] == 1]
fpr_male, tpr_male, _ = metrics.roc_curve(male['target'], male['score'])
```

```
female = marginals[marginals['gender'] == 0]
```

```


female = marginals[marginals['gender'] == 0]
fpr_female, tpr_female, _ = metrics.roc_curve(female['target'], female['score'])

plt.figure(1, figsize=(6, 5))
plt.title('ROC curves by grouping gender')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')

plt.plot(fpr_male, tpr_male, label='male', linestyle='-', color='b')
plt.plot(fpr_female, tpr_female, label='female', linestyle='--', color='g')

plt.legend()
plt.show()

```



```

rural = marginals[marginals['address'] == 1]
fpr_rural, tpr_rural, _ = metrics.roc_curve(rural['target'], rural['score'])

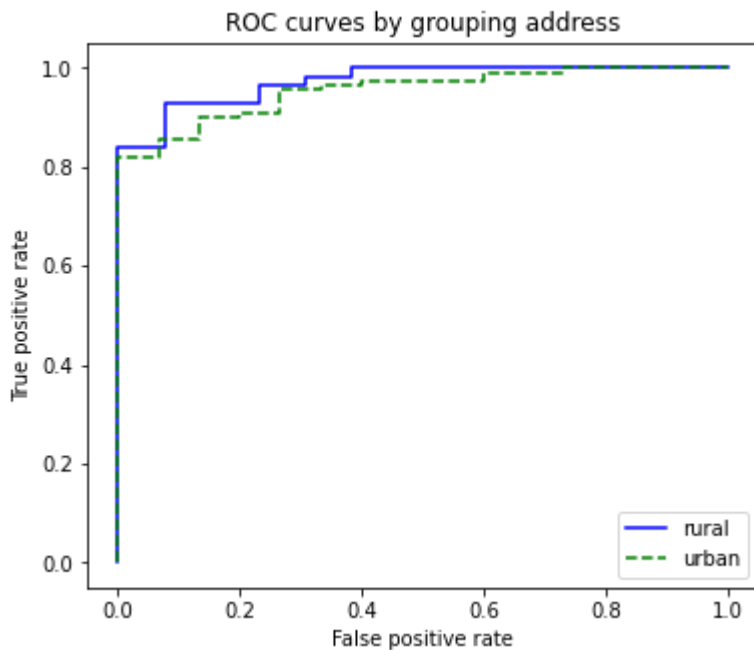
urban = marginals[marginals['address'] == 0]
fpr_urban, tpr_urban, _ = metrics.roc_curve(urban['target'], urban['score'])

plt.figure(1, figsize=(6, 5))
plt.title('ROC curves by grouping address')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')

plt.plot(fpr_rural, tpr_rural, label='rural', linestyle='-', color='b')
plt.plot(fpr_urban, tpr_urban, label='urban', linestyle='--', color='g')

plt.legend()
plt.show()

```



```

rural = marginals[marginals['address'] == 1]
fpr_rural, tpr_rural, _ = metrics.roc_curve(rural['target'], rural['score'])

ruralmale = rural[rural['gender'] == 1]
fpr_ruralmale, tpr_ruralmale, _ = metrics.roc_curve(ruralmale['target'], ruralmale['score'])

ruralfemale = rural[rural['gender'] == 0]
fpr_ruralfemale, tpr_ruralfemale, _ = metrics.roc_curve(ruralfemale['target'], ruralfemale['s

urban = marginals[marginals['address'] == 0]
fpr_urban, tpr_urban, _ = metrics.roc_curve(urban['target'], urban['score'])

urbanmale = urban[urban['gender'] == 1]
fpr_urbanmale, tpr_urbanmale, _ = metrics.roc_curve(urbanmale['target'], urbanmale['score'])

urbanfemale = urban[urban['gender'] == 0]
fpr_urbanfemale, tpr_urbanfemale, _ = metrics.roc_curve(urbanfemale['target'], urbanfemale['s

plt.figure(1, figsize=(6, 5))
plt.title('ROC curves by address and gender')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')

plt.plot(fpr_ruralmale, tpr_ruralmale, label='rural male', linestyle='-', color='b')
plt.plot(fpr_urbanmale, tpr_urbanmale, label='urban male', linestyle='--', color='g')
plt.plot(fpr_ruralfemale, tpr_ruralfemale, label='rural female', linestyle='-', color='r')
plt.plot(fpr_urbanfemale, tpr_urbanfemale, label='urban female', linestyle='--', color='y')

plt.legend()
plt.show()

```

